

```

In [1]: from sklearn.datasets import load_iris
        from sklearn.linear_model import LogisticRegression
        import numpy as np

        iris = load_iris()

In [4]: print(iris.data[0:5])

[[5.1 3.5 1.4 0.2]
 [4.9 3.  1.4 0.2]
 [4.7 3.2 1.3 0.2]
 [4.6 3.1 1.5 0.2]
 [5.  3.6 1.4 0.2]]

In [12]: X = iris.data
         y = iris.target

In [21]: X_new = (np.array([[3,4,5,2]]))

In [22]: #instantiate the model
         logreg = LogisticRegression()

         #fit the model with the data
         logreg.fit(X,y)

         #predict the response for new observations
         logreg.predict(X_new)

Out[22]: array([2])

In [24]: # Assessing LR Accuracy

         logreg = LogisticRegression()
         logreg.fit(X,y)
         # predict teh response values for the observations in X
         logreg.predict(X)

Out[24]: array([0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
                0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
                0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
                2, 1, 1, 1, 2, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 2, 2, 1, 1,
                1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 2, 2, 2, 2, 2, 2, 2, 2, 2,
                2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 1, 2, 2,
                2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2])

In [25]: # Storing the predict response values in y_pred
         y_pred = logreg.predict(X)

         # Checking length of y_pred
         len(y_pred)

Out[25]: 150

In [27]: # Calculating classification training accuracy using LR
         from sklearn import metrics
         print(metrics.accuracy_score(y, y_pred))

0.96

In [28]: # Using train_test_split for better accuracy evaluation

         from sklearn.model_selection import train_test_split
         X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=.4, random_state = 4)

In [29]: # training the model on the training set
         logreg = LogisticRegression()
         logreg.fit(X_train, y_train)

Out[29]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                             intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
                             penalty='l2', random_state=None, solver='liblinear', tol=0.0001,
                             verbose=0, warm_start=False)

In [30]: # make predictions on the testing set
         y_pred = logreg.predict(X_test)

         # comparing actual response values (y_test) with predicted response values (y_pred)
         print(metrics.accuracy_score(y_test,y_pred))

0.95

```